

Model Description

For Multi-class logistic regression, the softmax function is given as

$$p(y=k|x) = \sigma_m(f_k(x)) = \frac{e^{f_k(x)}}{\sum_{k=1}^K e^{f_k(x)}}$$

Where $f_k(x) = xw_k + w_{k0}$
 if $\theta = \begin{Bmatrix} w_k \\ w_{k0} \end{Bmatrix}$ and $x = [x \ 1]$

Hence then $f_k(x) = x\theta_k$

$$p(y=k|x) = \frac{\exp(x\theta_k)}{\sum_{k=1}^K \exp(x\theta_k)}$$

Using cross entropy loss function

$$\text{Loss} = - \sum_{m=1}^M \sum_{k=1}^K t[m][k] \log p(y[m]=k | x[m], \theta)$$

$$\text{Gradient} = - \sum_{m=1}^M \{t[m][k] - \sigma_m(x[m]\theta_k)\} x[m]$$

Where $x \in \mathbb{R}^{M \times N}$

$\theta \in \mathbb{R}^{N \times K}$

$t[m] = [0, 0, 1, 0, 0]$ if $x[m]$ belongs to class 3

Weight update:

$$\theta^{t+1} = \theta^t - \eta \times \text{Gradient}$$

Where η is the learning rate.

Experimental Settings

- Python Version 3.8 was used for coding.
- Training and testing data were imported and processed using the hints provided. The data were normalized by dividing with 255.0
- Training and testing labels were used to obtain $t[m]$
- Tensorflow Version 2.4.0 was used to build the softmax function and gradient and weight update
- Training was done by calling the classifier function and training and loss accuracies and loss were computed.
- Weights are saved.

- Testing data were separated into individual digits and their classification errors calculated. Overall errors are also calculated.
- Weights of individual classes and testing, training errors and accuracies plotted.

Hyper-parameters

learning rate = 0.75
iteration = 2000 (was decided from the results obtained using higher values).

Figures to attach here

- (1) training testing error curve
- (2) training testing accuracy curve
- (3) Weights (1 to 5)

Classification Errors

Digit 1 = 0.0141
Digit 2 = 0.0525
Digit 3 = 0.0680
Digit 4 = 0.0254
Digit 5 = 0.0776
average = 0.0468